


Control Flow

- Processors do only one thing:
 - From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
 - This sequence is the CPU's *control flow* (or *flow of control*)

Physical control flow




Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Exceptions

- An **exception** is a transfer of control to the OS *kernel* in response to some **event** (i.e., change in processor state)
 - Kernel is the memory-resident part of the OS
 - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Altering the Control Flow

- Up to now: two mechanisms for changing control flow:
 - Jumps and branches
 - Call and return

React to changes in **program state**

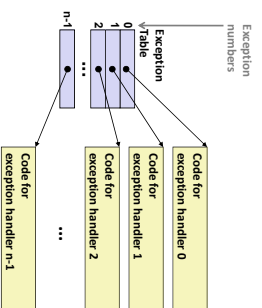
- Insufficient for a useful system:
 - Difficult to react to changes in **system state**
 - Data arrives from a disk or a network adapter
 - Instruction divides by zero
 - User hits Ctrl-C at the keyboard
 - System timer expires

- System needs mechanisms for “exceptional control flow”

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Exception Tables



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Exceptional Control Flow

- Exists at all levels of a computer system
- Low level mechanisms
 - Exceptions
 - Change in control flow in response to a system event (i.e., change in system state)
 - Implemented using combination of hardware and OS software
- Higher level mechanisms
 - Process context switch
 - Implemented by OS software and hardware timer
 - Signals
 - Implemented by OS software
 - Nonlocal jumps: setjmp () and longjmp ()
 - Implemented by C runtime library

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Asynchronous Exceptions (Interrupts)

- Caused by events external to the processor
 - Indicated by setting the processor's *interrupt pin*
 - Handler returns to “next” instruction
- Examples:
 - Timer interrupt
 - Every few ms, an external timer chip triggers an interrupt
 - Used by the kernel to take back control from user programs
 - I/O interrupt from external device
 - Hitting Ctrl-C at the keyboard
 - Arrival of a packet from a network
 - Arrival of data from a disk

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:
 - Traps**
 - Intentional
 - Examples: *system calls*, breakpoint traps, special instructions
 - Returns control to “next” instruction
 - Faults**
 - Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
 - Either re-executes faulting (“current”) instruction or aborts
 - Aborts**
 - Unintentional and unrecoverable
 - Examples: illegal instruction, parity error, machine check
 - Aborts current program

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



System Calls

- Each x86-64 system call has a unique ID number
- Examples:

Number	Nome	Description
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




System Call Example: Opening File

- User calls: `open (filename, options)`
- Calls `__open` function, which invokes system call instruction `syscall`

```

000000000006d70 <__open>:
...
e6d79: b8 02 00 00 00    mov $0x2,%eax # open is syscall #2
e6d7e: 0f 05             syscall # Return value in %rax
e6d80: 48 3d 01 01 ff ff  cmp $0xffffffff01,%rax
...
e6da: c3               retq
  
```



- `%rax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- Return value in `%rax`
- Negative value is an error corresponding to negative `errno`

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Fault Example: Page Fault


- User writes to memory location
- That portion (page) of user's memory is currently on disk

```

int a[1000];
main ()
{
    a[500] = 13;
}
  
```

```

80483b7: c7 05 10 94 04 08 0d movl $0xd,0x8049410
  
```



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




Fault Example: Invalid Memory Reference

```

int a[1000];
main ()
{
    a[5000] = 13;
}
  
```

```

80483b7: c7 05 60 a3 04 08 0d movl $0xd,0x804a360
  
```



- Sends `SIGSEGV` signal to user process
- User process exits with “segmentation fault”

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

