

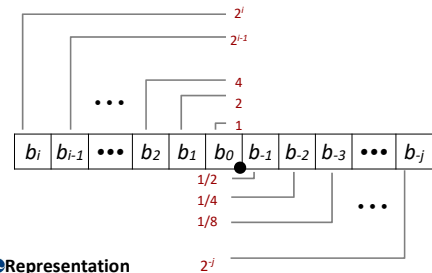
## Floating Point

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



1

## Fractional Binary Numbers



### Representation

- Bits to right of "binary point" represent fractional powers of 2
- Represents rational number: 
$$\sum_{k=-j}^i b_k \times 2^k$$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



4

## Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



2

## Fractional Binary Numbers: Examples

Value	Representation
5 3/4	101.11 <sub>2</sub>
2 7/8	10.111 <sub>2</sub>
1 7/16	1.0111 <sub>2</sub>

### Observations

- Divide by 2 by shifting right (unsigned)
- Multiply by 2 by shifting left
- Numbers of form 0.11111...<sub>2</sub> are just below 1.0
  - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
  - Use notation  $1.0 - \epsilon$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



5

## Fractional binary numbers

- What is 1011.101<sub>2</sub>?

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



3

## Representable Numbers

### Limitation #1

- Can only exactly represent numbers of the form  $x/2^k$ 
  - Other rational numbers have repeating bit representations
- Value Representation
- 1/3 0.01010101 [01]...<sub>2</sub>
- 1/5 0.001100110011 [0011]...<sub>2</sub>
- 1/10 0.0001100110011 [0011]...<sub>2</sub>

### Limitation #2

- Just one setting of binary point within the w bits
  - Limited range of numbers (very small values? very large?)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



6

## Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




7

## Precision options


### Single precision: 32 bits



### Double precision: 64 bits



### Extended precision: 80 bits (Intel only)



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



10

## IEEE Floating Point

- IEEE Standard 754
  - Established in 1985 as uniform standard for floating point arithmetic
    - Before that, many idiosyncratic formats
  - Supported by all major CPUs
- Driven by numerical concerns
  - Nice standards for rounding, overflow, underflow
  - Hard to make fast in hardware
    - Numerical analysts predominated over hardware designers in defining standard

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



8

## "Normalized" Values

$$v = (-1)^s M 2^E$$

### When: $\text{exp} \neq 000\dots 0$ and $\text{exp} \neq 111\dots 1$

#### Exponent coded as a *biased* value: $E = \text{Exp} - \text{Bias}$

- Exp: unsigned value of exp field
- Bias =  $2^{k-1} - 1$ , where  $k$  is number of exponent bits
  - Single precision: 127 (Exp: 1...254, E: -126...127)
  - Double precision: 1023 (Exp: 1...2046, E: -1022...1023)

#### Significand coded with implied leading 1: $M = 1.\text{xxx}\dots\text{x}_2$

- xxx...x: bits of frac field
- Minimum when  $\text{frac} = 000\dots 0$  ( $M = 1.0$ )
- Maximum when  $\text{frac} = 111\dots 1$  ( $M = 2.0 - \epsilon$ )
- Get extra leading bit for "free"

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



11

## Floating Point Representation

### Numerical Form:

$$(-1)^s M 2^E$$

- Sign bit  $s$  determines whether number is negative or positive
- Significand  $M$  normally a fractional value in range  $[1.0, 2.0)$ .
- Exponent  $E$  weights value by power of two

### Encoding

- MSB  $s$  is sign bit
- exp field encodes  $E$  (but is not equal to  $E$ )
- frac field encodes  $M$  (but is not equal to  $M$ )



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



9

## Normalized Encoding Example

$$v = (-1)^s M 2^E$$
$$E = \text{Exp} - \text{Bias}$$

### Value: float $F = 15213.0$ ;

$$15213_{10} = 11101101101101_2$$
$$= 1.1101101101101_2 \times 2^{13}$$


### Significand

$$M = 1.1101101101101_2$$
$$\text{frac} = 1101101101101000000000_2$$

### Exponent

$$E = 13$$
$$\text{Bias} = 127$$
$$\text{Exp} = 140 = 10001100_2$$

### Result:



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



12

## Denormalized Values

$$V = (-1)^s M 2^E$$

$$E = 1 - \text{Bias}$$

- Condition:  $\text{exp} = 000\dots 0$
- Exponent value:  $E = 1 - \text{Bias}$  (instead of  $E = 0 - \text{Bias}$ )
- Significand coded with implied leading 0:  $M = 0.\text{xxx}\dots\text{x}_2$ 
  - $\text{xxx}\dots\text{x}$ : bits of  $\text{frac}$
- Cases
  - $\text{exp} = 000\dots 0, \text{frac} = 000\dots 0$ 
    - Represents zero value
    - Note distinct values: +0 and -0
  - $\text{exp} = 000\dots 0, \text{frac} \neq 000\dots 0$ 
    - Numbers closest to 0.0
    - Equispaced

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



13

## Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



16

## Special Values

- Condition:  $\text{exp} = 111\dots 1$
- Case:  $\text{exp} = 111\dots 1, \text{frac} = 000\dots 0$ 
  - Represents value  $\infty$  (infinity)
  - Operation that overflows
  - Both positive and negative
  - E.g.,  $1.0/0.0 = -1.0/-0.0 = +\infty$ ,  $1.0/-0.0 = -\infty$
- Case:  $\text{exp} = 111\dots 1, \text{frac} \neq 000\dots 0$ 
  - Not-a-Number (NaN)
  - Represents case when no numeric value can be determined
  - E.g.,  $\text{sqrt}(-1)$ ,  $\infty - \infty$ ,  $\infty \times 0$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



14

## Tiny Floating Point Example




- 8-bit Floating Point Representation
  - the sign bit is in the most significant bit
  - the next four bits are the exponent, with a bias of 7
  - the last three bits are the  $\text{frac}$
- Same general form as IEEE Format
  - normalized, denormalized
  - representation of 0, NaN, infinity

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



17

## Visualization: Floating Point Encodings



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



15

## Dynamic Range (Positive Only)

					$V = (-1)^s M 2^E$ $n: E = \text{Exp} - \text{Bias}$ $d: E = 1 - \text{Bias}$	
	s	exp	frac	E	Value	
Denormalized numbers	0	0000	000	-6	0	closest to zero
	0	0000	001	-6	$1/8 \times 1/64 = 1/512$	
	0	0000	010	-6	$2/8 \times 1/64 = 2/512$	
	...					
	0	0000	110	-6	$6/8 \times 1/64 = 6/512$	largest denorm
Normalized numbers	0	0000	111	-6	$7/8 \times 1/64 = 7/512$	
	0	0001	000	-6	$8/8 \times 1/64 = 8/512$	smallest norm
	0	0001	001	-6	$9/8 \times 1/64 = 9/512$	
	...					
	0	0110	110	-1	$14/8 \times 1/2 = 14/16$	closest to 1 below
	0	0110	111	-1	$15/8 \times 1/2 = 15/16$	
	0	0111	000	0	$8/8 \times 1 = 1$	closest to 1 above
	0	0111	001	0	$9/8 \times 1 = 9/8$	
	0	0111	010	0	$10/8 \times 1 = 10/8$	largest norm
	...					
	0	1110	110	7	$14/8 \times 128 = 224$	
	0	1110	111	7	$15/8 \times 128 = 240$	
	0	1111	000	n/a	inf	

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition




18


## Distribution of Values

### 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- Bias is  $2^{3-1}-1 = 3$



### Notice how the distribution gets denser toward zero.



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



19

## Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication**
- Floating point in C
- Summary

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition





22

## Distribution of Values (close-up view)

### 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- Bias is 3



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



20

## Floating Point Operations: Basic Idea

$$x \oplus_f y = \text{Round}(x \oplus y)$$

$$x \otimes_f y = \text{Round}(x \otimes y)$$

### Basic idea

- First **compute exact result**
- Make it fit into desired precision
  - Possibly overflow if exponent too large
  - Possibly **round to fit into frac**

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



23

## Special Properties of the IEEE Encoding

### FP Zero Same as Integer Zero

- All bits = 0

### Can (Almost) Use Unsigned Integer Comparison

- Must first compare sign bits
- Must consider  $-0 = 0$
- NaNs problematic
  - Will be greater than any other values
  - What should comparison yield?
- Otherwise OK
  - Denorm vs. normalized
  - Normalized vs. infinity

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



21

## Rounding

### Rounding Modes (illustrate with \$ rounding)

	\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
Towards zero	\$1	\$1	\$1	\$2	-\$1
Round down ( $-\infty$ )	\$1	\$1	\$1	\$2	-\$2
Round up ( $+\infty$ )	\$2	\$2	\$2	\$3	-\$1
Nearest Even (default)	\$1	\$2	\$2	\$2	-\$2

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



24

## Closer Look at Round-To-Even

### Default Rounding Mode

- Hard to get any other kind without dropping into assembly
- All others are statistically biased
  - Sum of set of positive numbers will consistently be over- or under-estimated

### Applying to Other Decimal Places / Bit Positions

- When exactly halfway between two possible values
  - Round so that least significant digit is even
- E.g., round to nearest hundredth
 

7.8949999	7.89	(Less than half way)
7.8950001	7.90	(Greater than half way)
7.8950000	7.90	(Half way—round up)
7.8850000	7.88	(Half way—round down)

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

25

## Floating Point Addition


$$(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$$

- Assume  $E1 > E2$

Get binary points lined up

### Exact Result: $(-1)^s M 2^E$

- Sign  $s$ , significand  $M$ :
  - Result of signed align & add
- Exponent  $E$ :  $E1$



### Fixing

- If  $M \geq 2$ , shift  $M$  right, increment  $E$
- If  $M < 1$ , shift  $M$  left  $k$  positions, decrement  $E$  by  $k$
- Overflow if  $E$  out of range
- Round  $M$  to fit **frac** precision

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

28

## Rounding Binary Numbers

### Binary Fractional Numbers

- "Even" when least significant bit is 0
- "Half way" when bits to right of rounding position = 100...2

### Examples

- Round to nearest 1/4 (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded
2 3/32	10.000112	10.002	(<1/2—down)	2
2 3/16	10.001102	10.012	(>1/2—up)	2 1/4
2 7/8	10.111002	11.002	(1/2—up)	3
2 5/8	10.101002	10.102	(1/2—down)	2 1/2

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

26

## Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

29

## FP Multiplication

$$(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$$

### Exact Result: $(-1)^s M 2^E$

- Sign  $s$ :  $s1 \wedge s2$
- Significand  $M$ :  $M1 \times M2$
- Exponent  $E$ :  $E1 + E2$

### Fixing

- If  $M \geq 2$ , shift  $M$  right, increment  $E$
- If  $E$  out of range, overflow
- Round  $M$  to fit **frac** precision

### Implementation

- Biggest chore is multiplying significands

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

27

## Floating Point in C

### C Guarantees Two Levels

- float** single precision
- double** double precision

### Conversions/Casting

- Casting between **int**, **float**, and **double** changes bit representation
- double/float → int**
  - Truncates fractional part
  - Like rounding toward zero
  - Not defined when out of range or NaN: Generally sets to TMin
- int → double**
  - Exact conversion, as long as **int** has ≤ 53 bit word size
- int → float**
  - Will round according to rounding mode

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

30

## Floating Point Puzzles

🔊 For each of the following C expressions, either:

- 🔊 Argue that it is true for all argument values
- 🔊 Explain why not true

```
int x = ...;
float f = ...;
double d = ...;
```

Assume neither  
d nor f is NaN

- `x == (int)(float) x`
- `x == (int)(double) x`
- `f == (float)(double) f`
- `d == (double)(float) d`
- `f == -(-f);`
- `2/3 == 2/3.0`
- `d < 0.0`       $\Rightarrow$     `((d*2) < 0.0)`
- `d > f`       $\Rightarrow$     `-f > -d`
- `d * d >= 0.0`
- `(d+f)-d == f`



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

31

## Summary

- 🔊 IEEE Floating Point has clear mathematical properties
- 🔊 Represents numbers of form  $M \times 2^E$
- 🔊 One can reason about operations independent of implementation
  - 🔊 As if computed with perfect precision and then rounded
- 🔊 Not the same as real arithmetic
  - 🔊 Violates associativity/distributivity
  - 🔊 Makes life difficult for compilers & serious numerical applications programmers



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

32