

Today

- Linking
- Case study: Library interpositioning

Byrni and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Why Linkers?

- Reason 1: Modularity
 - Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library

Byrni and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Example C Program

```
int sum(int* a, int n);
int array[2] = {1, 2};
int main()
{
    int val = sum(array, 2);
    return val;
}
main.c
```

```
int sum(int* a, int n)
{
    int i, s = 0;
    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}
sum.c
```

Byrni and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Why Linkers? (cont)


- Reason 2: Efficiency
 - Time: Separate compilation
 - Change one source file, compile, and then relink.
 - No need to recompile other source files.
- Space: Libraries
 - Common functions can be aggregated into a single file...
 - Yet executable files and running memory images contain only code for the functions they actually use.

Byrni and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Static Linking

- Programs are translated and linked using a compiler driver.
 - `linux> gcc -Og -o prog main.c sum.c`
 - `linux> ./prog`



Byrni and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



What Do Linkers Do?

- Step 1: Symbol resolution
 - Programs define and reference symbols (global variables and functions):
 - `void swap() {...} /* define symbol swap */`
 - `swap(); /* reference symbol swap */`
 - `int *xp = &x; /* define symbol xp, reference x */`
 - Symbol definitions are stored in object file (by assembler) in symbol table.
 - Symbol table is an array of structs
 - Each entry includes name, size, and location of symbol.
- During symbol resolution step, the linker associates each symbol reference with exactly one symbol definition.

Byrni and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



What Do Linkers Do? (cont)

■ Step 2: Relocation

- Merges separate code and data sections into single sections
- Relocates symbols from their relative locations in the `.o` files to their final absolute memory locations in the executable.
- Updates all references to these symbols to reflect their new positions.

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Linker Symbols

■ Global symbols

- Symbols defined by module `m` that can be referenced by other modules.
- E.g.: `non-static` C functions and `non-static` global variables.

■ External symbols

- Global symbols that are referenced by module `m` but defined by some other module.

■ Local symbols

- Symbols that are defined and referenced exclusively by module `m`.
- E.g.: C functions and global variables defined with the `static` attribute.

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Three Kinds of Object Files (Modules)

■ Relocatable object file (`.o` file)

- Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each `.o` file is produced from exactly one source `.c` file

■ Executable object file (`a.out` file)

- Contains code and data in a form that can be copied directly into memory and then executed.


■ Shared object file (`.so` file)

- Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
- Called *Dynamic Link Libraries* (DLLs) by Windows

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Step 1: Symbol Resolution



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Executable and Linkable Format (ELF)

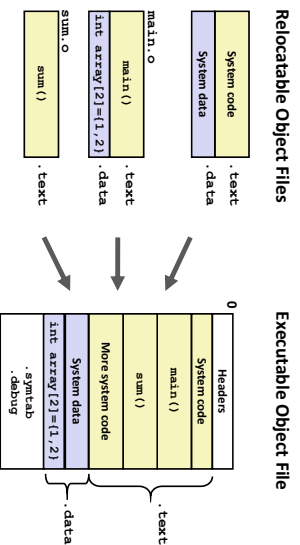
■ Standard binary format for object files

- **One unified format for**
 - Relocatable object files (`.o`)
 - Executable object files (`a.out`)
 - Shared object files (`.so`)
- **Generic name: ELF binaries**

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Step 2: Relocation



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Packaging Commonly Used Functions

- **How to package functions commonly used by programmers?**
 - Math, I/O, memory management, string manipulation, etc.
- **Awkward, given the linker framework so far:**
 - **Option 1:** Put all functions into a single source file
 - Programmers link big object file into their programs
 - Space and time inefficient
 - **Option 2:** Put each function in a separate source file
 - Programmers explicitly link appropriate binaries into their programs
 - More efficient, but burdensome on the programmer

Byrns and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



22

Commonly Used Libraries

- **libc.a (the C standard library)**
 - 4.6 MB archive of 1496 object files.
 - I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math
- **libm.a (the C math library)**
 - 2 MB archive of 444 object files.
 - floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t libc.a | sort
...
fork.o
...
fprintf.o
fpconvctrl.o
fpconv.o
freopen.o
fscanf.o
fseek.o
ftab.o
```

```
% ar -t libm.a | sort
...
__acos.o
__acosh.o
__acoshl.o
__acoshl.o
__acoshl.o
__asin.o
__asinh.o
__asinh.o
__asinh.o
```

Byrns and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



23

Old-fashioned Solution: Static Libraries

- **Static libraries (a archive files)**
 - Concatenate related relocatable object files into a single file with an index (called an *archive*).
 - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
 - If an archive member file resolves reference, link it into the executable.

Byrns and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



24

Linking with Static Libraries

```
#include <stdio.h>
#include "vector.h"

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main()
{
    addvec(x, y, z, 2);
    printf("z = [%d %d]\n",
           z[0], z[1]);
    return 0;
}

main2.c
```

```
void addvec(int *x, int *y,
            int *z, int n) {
    int i;
    for (i = 0; i < n; i++)
        z[i] = x[i] + y[i];
}

addvec.o

void multvec(int *x, int *y,
             int *z, int n)
{
    int i;
    for (i = 0; i < n; i++)
        z[i] = x[i] * y[i];
}


multvec.o
```

Byrns and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



27

Creating Static Libraries



C standard library

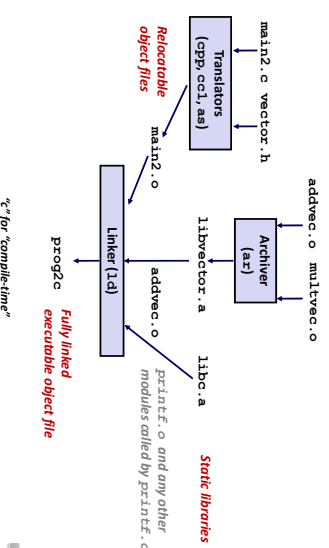
- Archiver allows incremental updates
- Recompile function that changes and replace .o file in archive.

Byrns and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



25

Linking with Static Libraries



"C" for "compile-time"

Fully linked executable object file

Byrns and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



28

Using Static Libraries

- **Linker's algorithm for resolving external references:**
 - Scan, **.o** files and **.a** files in the command line order.
 - During the scan, keep a list of the current unresolved references.
 - As each new **.o** or **.a** file, **obj**, is encountered, try to resolve each unresolved reference in the list against the symbols defined in **obj**. If any entries in the unresolved list at and of scan, then error.

■ **Problem:**

- Command line order matters!
- Moral: put libraries at the end of the command line.

```
unix> gcc -L. libtest.o -lmine
unix> gcc -L. -lmine libtest.o
libtest.o: In function 'main':
```



Modern Solution: Shared Libraries

- **Static libraries have the following disadvantages:**
 - Duplication in the stored executables (every function needs libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to explicitly relink
- **Modern solution: Shared Libraries**
 - Object files that contain code and data that are loaded and linked into an application *dynamically*, at either *load-time* or *run-time*
 - Also called: dynamic link libraries, DLLs, *.so* files

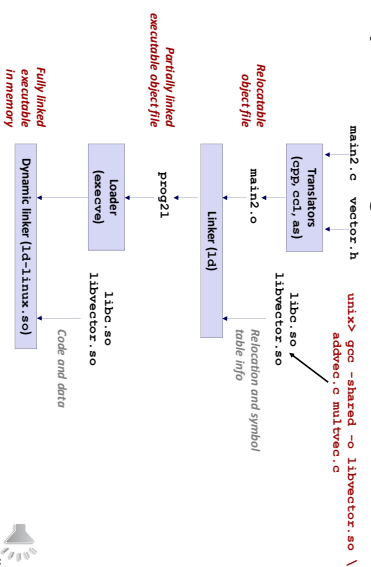


Shared Libraries (cont.)

- **Dynamic linking can occur when executable is first loaded and run (load-time linking).**
 - Common case for Linux, handled automatically by the dynamic linker (`ld-linux.so`).
 - Standard C library (`libc.so`) usually dynamically linked.
- **Dynamic linking can also occur after program has begun (run-time linking).**
 - In Linux, this is done by calls to the `dlopen()` interface.
 - Distributing software.
 - High-performance web servers.
 - Runtime library repositioning.
- **Shared library routines can be shared by multiple processes**
 - More on this when we learn about virtual memory



Dynamic Linking at Load-time



Linking Summary

- **Linking** is a technique that allows programs to be constructed from multiple object files.
- **Linking can happen at different times in a program's lifetime:**
 - Compile time (when a program is compiled)
 - Load time (when a program is loaded into memory)
 - Run time (while a program is executing)
- **Understanding linking can help you avoid nasty errors and make you a better programmer.**



Today

- **Linking**
- **Case study: Library interpositioning**



Case Study: Library Interpositioning

- **Library interpositioning** : powerful linking technique that allows programmers to intercept calls to arbitrary functions
- **Interpositioning can occur at:**
 - Compile time: When the source code is compiled
 - Link time: When the relocatable object files are statically linked to form an executable object file
 - Load/run time: When an executable object file is loaded into memory, dynamically linked, and then executed.

Byrns and Orlikowski, Computer Systems: A Programmer's Perspective, Third Edition



37

Some Interpositioning Applications

- **Security**
 - Confinement (sandboxing)
 - Behind the scenes encryption
- **Debugging**
 - In 2014, two Facebook engineers debugged a treacherous 1-year old bug in their iPhone app using interpositioning
 - Code in the SPDY networking stack was writing to the wrong location
 - Solved by intercepting calls to Posix write functions (write, writew, pwrite)

Source: Facebook engineering blog post at <https://code.facebook.com/post/3130347212144/debugging-file-corruption-on-ios/>

Byrns and Orlikowski, Computer Systems: A Programmer's Perspective, Third Edition



38

Some Interpositioning Applications

- **Monitoring and Profiling**
 - Count number of calls to functions
 - Characterize call sites and arguments to functions
- Malloc tracing
 - Detecting memory leaks
- **Generating address traces**

Byrns and Orlikowski, Computer Systems: A Programmer's Perspective, Third Edition



39