

# Machine-Level Programming V: Advanced Topics



## Memory Allocation Example


not drawn to scale

```
char big_array[1L<<24]; /* 16 MB */
char huge_array[1L<<31]; /* 2 GB */

int global = 0;

int useless() { return 0; }

int main()
{
    void *p1, *p2, *p3, *p4;
    int local = 0;
    p1 = malloc(1L << 28); /* 256 MB */
    p2 = malloc(1L << 8); /* 256 B */
    p3 = malloc(1L << 32); /* 4 GB */
    p4 = malloc(1L << 8); /* 256 B */
    /* Some print statements ... */
}
```



## Today

- Memory Layout
- Buffer Overflow
- Vulnerability
- Protection
- Unions




## x86-64 Example Addresses

address range ~2<sup>47</sup>

```
Local
p1
p3
p4
p2
big_array
huge_array
main()
useless()
```

```
0x00007f2e4d3ba87c
0x00007f7162a1d010
0x0000000083594d20
0x0000000083594d10
0x0000000080601060
0x0000000000601060
0x000000000040060c
0x0000000000400590
```




## x86-64 Linux Memory Layout

not drawn to scale

- Stack
- Runtime stack (8MB limit)
- E.g., local variables
- Heap
- Dynamically allocated as needed
- When call malloc(), calloc(), new()
- Data
- Statically allocated data
- E.g., global vars, static vars, string constants
- Text / Shared Libraries
- Executable machine instructions
- Read-only

Hex Address 4000000 0000000



## Today

- Memory Layout
- Buffer Overflow
- Vulnerability
- Protection
- Unions



## Recall: Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    return s.d;
}
```

```
fun(0) 0x 3.14
fun(1) 0x 3.14
fun(2) 0x 3.139998664856
fun(3) 0x 2.00000061035156
fun(4) 0x 3.14
fun(6) 0x Segmentation fault
```

Byrns and Chellison, Computer Systems: A Programmer's Perspective, Third Edition



## String Library Code

### Implementation of Unix function gets()

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- No way to specify limit on number of characters to read
- Similar problems with other library functions
- strcpy, strcat: Copy strings of arbitrary length
- scanf, fscanf, sscanf, when given %s conversion specification

Byrns and Chellison, Computer Systems: A Programmer's Perspective, Third Edition




## Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

fun(0) 0x 3.14
fun(1) 0x 3.14
fun(2) 0x 3.139998664856
fun(3) 0x 2.00000061035156
fun(4) 0x 3.14
fun(6) 0x Segmentation fault
```

### Explanation:



Byrns and Chellison, Computer Systems: A Programmer's Perspective, Third Edition



## Vulnerable Buffer Code

```
/* Echo line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

← btw, how big is buf enough?

```
unix> ./bufdemo-nop
Type a string: 012345678901234567890123
012345678901234567890123

unix> ./bufdemo-nop
Type a string: 0123456789012345678901234
Segmentation Fault
```

Byrns and Chellison, Computer Systems: A Programmer's Perspective, Third Edition



## Such problems are a BIG deal

- Generally called a "buffer overflow"
- when exceeding the memory size allocated for an array
- Why a big deal?
  - it's the #1 technical cause of security vulnerabilities
  - #1 overall cause is social engineering / user ignorance
- Most common form
  - Unchecked lengths on string inputs
  - Particularly for bounded character arrays on the stack
  - sometimes referred to as stack smashing

Byrns and Chellison, Computer Systems: A Programmer's Perspective, Third Edition



## Buffer Overflow Disassembly

echo:

```
0000000004006cf: 0000000004006cf <echo>:
4006cf: 48 83 ec 18      sub    $0x18,%rax
4006d3: 48 89 e7        mov    %rax,%rdi
4006d6: e8 a5 ff ff ff  callq 400680 <gets>
4006db: 48 89 e7        mov    %rax,%rdi
4006de: e8 3d fe ff ff  callq 400520 <puts@plt>
4006e3: 48 83 c4 18      add    $0x18,%rax
4006e7: c3             retq
```

```
call echo:
4006e8: 48 83 ec 08      sub    $0x8,%rax
4006ec: b8 00 00 00 00  mov    $0x0,%eax
4006f1: e8 d9 ff ff ff  callq 4006cf <echo>
4006f6: 48 83 c4 08      add    $0x8,%rax
4006fa: c3             retq
```

Byrns and Chellison, Computer Systems: A Programmer's Perspective, Third Edition



## Buffer Overflow Stack

*Before call to gets*

Stack Frame for call_echo	
Return Address (8 bytes)	
20 bytes unused	
[3][2][1][0] buf ← %rsp	

```
/* Echo line */
void echo ()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

echo:	subq \$24, %rsp	movq %rsp, %rdi	call gets	...
-------	-----------------	-----------------	-----------	-----

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



## Buffer Overflow Stack Example #2

*After call to gets*

Stack Frame for call_echo	
00 00 00 00	
00 40 00 <b>34</b>	
33 32 31 30	
39 38 37 36	
35 34 33 32	
31 30 39 38	
37 36 35 34	
33 32 31 30	

```
void echo ()
{
    char buf[4];
    gets(buf);
    ...
}
```

```
call_echo:
    callq 4006cf <echo>
    add $0x8, %rsp
    ...
```

buf ← %rsp
------------

```
unix> ./bufdemo-nsp
Type a string: 0123456789012345678901234
Segmentation fault
```

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Overflowed buffer and corrupted return pointer

## Buffer Overflow Stack Example

*Before call to gets*

Stack Frame for call_echo	
00 00 00 00	
00 40 06 f6	
20 bytes unused	
[3][2][1][0] buf ← %rsp	

```
void echo ()
{
    char buf[4];
    gets(buf);
    ...
}
```

```
call_echo:
    callq 4006cf <echo>
    add $0x8, %rsp
    ...
```

echo:	subq \$24, %rsp	movq %rsp, %rdi	call gets	...
-------	-----------------	-----------------	-----------	-----

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



## Buffer Overflow Stack Example #3

*After call to gets*

Stack Frame for call_echo	
00 00 00 00	
00 40 06 00	
33 32 31 30	
39 38 37 36	
35 34 33 32	
31 30 39 38	
37 36 35 34	
33 32 31 30	

```
void echo ()
{
    char buf[4];
    gets(buf);
    ...
}
```

```
call_echo:
    callq 4006cf <echo>
    add $0x8, %rsp
    ...
```

buf ← %rsp
------------

```
unix> ./bufdemo-nsp
Type a string: 012345678901234567890123
012345678901234567890123
```

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Overflowed buffer, corrupted return pointer, but program seems to work!

## Buffer Overflow Stack Example #1

*After call to gets*

Stack Frame for call_echo	
00 00 00 00	
00 40 06 f6	
<b>00 32 31 30</b>	
39 38 37 36	
35 34 33 32	
31 30 39 38	
37 36 35 34	
33 32 31 30	

```
void echo ()
{
    char buf[4];
    gets(buf);
    ...
}
```

```
call_echo:
    callq 4006cf <echo>
    add $0x8, %rsp
    ...
```

buf ← %rsp
------------

```
unix> ./bufdemo-nsp
Type a string: 01234567890123456789012
01234567890123456789012
```

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Overflowed buffer, but did not corrupt state

## Buffer Overflow Stack Example #3 Explained

*After call to gets*

Stack Frame for call_echo	
00 00 00 00	
00 40 06 <b>00</b>	
<b>33 32 31 30</b>	
39 38 37 36	
35 34 33 32	
31 30 39 38	
37 36 35 34	
33 32 31 30	

```
register tm clones:
    mov %rsp, %rdp
    mov %rax, %rdx
    shr $0x3f, %rdx
    add 400609a, %rdi
    sar %rcx, %rax
    jne 400610, %rcx
    pop %rdp
    retq 400613
```

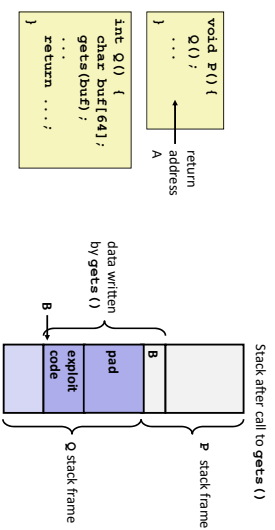
buf ← %rsp
------------

"Return" to unrelated code  
Lots of things happen, without modifying critical state  
Eventually executes retq back to main

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



## Code Injection Attacks



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer B
- When Q executes `ret`, will jump to exploit code

Bryant and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## Exploits Based on Buffer Overflows

- Buffer overflow bugs can allow remote machines to execute arbitrary code on victim machines**
- Distressingly common in real programs
- Programmers keep making the same mistakes ☹
- Recent measures make these attacks much more difficult
- Examples across the decades**
  - Original "Internet worm" (1988)
  - "IM wars" (1999)
  - Twilight hack on Wii (2000s)
  - ... and many, many more
- You will learn some of the tricks in attacklab**
- Hopefully to convince you to never leave such holes in your programs!

Bryant and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## Example: the original Internet worm (1988)


- Exploited a few vulnerabilities to spread**
  - Early versions of the finger server (`fingerd`) used `gets()` to read the argument sent by the client:
    - `finger drohles.cmu.edu`
  - Worm attacked fingerd server by sending phony argument:
    - `finger "exploit-code padding new-return-address"`
  - exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.
- Once on a machine, scanned for other machines to attack**
  - invaded ~6000 computers in hours (10% of the internet ☹)
  - see June 1989 article in *Comm. of the ACM*
  - the young author of the worm was prosecuted...
  - and CERT was formed... still homed at CMU

Bryant and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## Example 2: IM War

- July, 1999**
  - Microsoft launches MSN Messenger (instant messaging system).
  - Messenger clients can access popular AOL Instant Messaging Service (AIM) servers



Bryant and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## IM War (cont.)

- August 1999**
  - Mysteriously, Messenger clients can no longer access AIM servers
  - Microsoft and AOL begin the IM war:
    - AOL changes server to disallow Messenger clients
    - Microsoft makes changes to clients to defeat AOL changes
    - At least 13 such skirmishes
  - What was really happening?
    - AOL had discovered a buffer overflow bug in their own AIM clients
    - They exploited it to detect and block Microsoft: the exploit code returned a 4-byte signature (the bytes at some location in the AIM client) to server
    - When Microsoft changed code to match signature, AOL changed signature location

Bryant and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## OK, what to do about buffer overflow attacks

- Avoid overflow vulnerabilities
- Employ system-level protections
- Have compiler use "stack canaries"
- Lets talk about each...

Bryant and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## 1. Avoid Overflow Vulnerabilities in Code (!)

```
/* Echo line */
void echo()
{
    char buf[4]; /* Way too small! */
    fgets(buf, 4, stdin);
    puts(buf);
}
```

- For example, use library routines that limit string lengths
- `fgets` instead of `gets`
- `strncpy` instead of `strcpy`
- Don't use `scanf` with `%s` conversion specification
- Use `fgets` to read the string
- Or use `%ns` where `n` is a suitable integer

Byrnes and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## 3. Stack Canaries can help

- Idea
  - Place special value ("canary") on stack just beyond buffer
  - Check for corruption before exiting function
- GCC implementation
  - `-fstack-protector`
  - Now the default (disabled earlier)

```
unix> ./bufdemo-sp
Type a string: 0123456
0123456
```

```
unix> ./bufdemo-sp
Type a string: 01234567
*** stack smashing detected ***
```

Byrnes and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## 2. System-Level Protections can help

- Randomized stack offsets
  - At start of program, allocate random amount of space on stack
  - Shifts stack addresses for entire program
  - Makes it difficult for hacker to predict beginning of inserted code
  - E.g.: 5 executions of memory allocation code
- Stack repositioned each time program executes

```
local 0x7f643a8b7c 0x7f7f5a4f9c 0x7f6ea2780c 0x7f6e427d4c 0x7f6d42077c
B7 →
```



## Protected Buffer Disassembly

echo:


```
40072f: sub    $0x18,%rsp
400733: mov    %fs:0x28,%rax
40073c: mov    %rax,0x8(%rsp)
400741: xor    %eax,%eax
400743: mov    %rsp,%edi
400746: callq 4006a0<gets>
40074b: mov    %rsp,%edi
40074e: callq 400570<puts@plt>
400753: mov    0x8(%rsp),%rax
400758: xor    %fs:0x28,%rax
400761: je     400768<echo+0x39>
400763: callq 400580<__stack_chk_fail@plt>
40076c: retq   $0x18,%rsp
```

Byrnes and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## 2. System-Level Protections can help

- Nonexecutable code segments
  - In traditional x86, can mark "read-only" or "writable"
  - Can execute anything
  - readable
  - X86-64 added explicit "execute" permission
  - Stack marked as non-executable



Any attempt to execute this code will fail


Byrnes and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## Setting Up Canary

Before call to `gets`

```
/* Echo line */
void echo()
{
    char buf[4]; /* Way too small! */
    fgets(buf);
    puts(buf);
}
```



buf ← %rsp

```
echo:
    . . .
    movq %fs:40,%rax # Get canary
    movq %rax,8(%rsp) # Place on stack
    xorl %eax,%eax # Erase canary
    . . .
```

Byrnes and Chhabra, Computer Systems: A Programmer's Perspective, Third Edition



## Checking Canary

*After call to gets*

```
/* Echo line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

Stack frame for call_echo
Return Address (8 bytes)
Canary (8 bytes)
00 36 35 34 33 32 31 30

Input: 0123456  
buf ← %rsp

```
echo:
    .L6:
    movq    8(%rsp), %rax    # Retrieve from
stack      %fs:40, %rax    # Compare to canary
    xorq    %rsi, %rax      # If same, OK
    je      .L6             # If same, OK
    call    __stack_chk_fail # Fail
    .L6:
```

## Gadget Example #2

```
void setval(unsigned *p) {
    *p = 334763060u;
}
```


<setval>: 4004d9: c7 07 d4 48 89 c7 movl \$0xc78948d4, (%rdi)  
4004df: c3 retq  
rdi ← rax  
Gadget address = 0x4004dc

### Repurpose byte codes

## Return-Oriented Programming Attacks

- Challenge (for hackers)
  - Stack randomization makes it hard to predict buffer location
  - Marking stack nonexecutable makes it hard to insert binary code
- Alternative Strategy
  - Use existing code
    - E.g., library code from stdlib
  - String together fragments to achieve overall desired outcome
  - Does not overcome stack canaries
- Construct program from gadgets
  - Sequence of instructions ending in `ret`
    - Encoded by single byte `0xc3`
  - Code positions fixed from run to run
  - Code is executable

## ROP Execution



- Trigger with `ret` instruction
  - Will start executing Gadget 1
- Final `ret` in each gadget will start next one

## Gadget Example #1

```
long ab_plus_c
(long a, long b, long c)
{
    return a*b + c;
}
```

00000000004004d0 <ab\_plus\_c>:  
4004d0: 48 0f af f6 imul %rsi, %rdi  
4004d4: 48 8d 04 17 lea (%rdi, %rcx, 1), %rax  
4004d8: c3 retq

rax ← rdi + rdx  
Gadget address = 0x4004d4


### Use tail end of existing functions

## Today

- Memory Layout
- Buffer Overflow
- Vulnerability
- Protection
- Unions

# Union Allocation

- Allocate according to largest element
- Can only use one field at a time



Bjorn and Christen, Computer Systems: A Programmer's Perspective, Third Edition

