


Array Allocation

Basic Principle


- $T \ A[L];$
- Array of data type T and length L
- Contiguously allocated region of $L * \text{sizeof}(T)$ bytes in memory



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

3

Array Accessing Example



```
int get_digit
(zip_dig z, int digit)
{
    return z[digit];
}
```

X86-64

```
# %rdi = z
# %rsi = digit
movl (%rdi,%rsi,4), %eax # z[digit]
```

- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at `%rdi + 4*%rsi`
- Use memory reference `(%rdi,%rsi,4)`

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

4

Today

Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

Structures

- Allocation
- Access
- Alignment


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

2

Array Example

```
#define ZLEN 5
typedef int zip_dig[ZLEN];

zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```



- Declaration `"zip_dig cmu"` equivalent to `"int cmu[5]"`
- Example arrays were allocated in successive 20 byte blocks
- Not guaranteed to happen in general

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


5

Machine-Level Programming IV: Data

Array Access

Basic Principle

- $T \ A[L];$
- Array of data type T and length L
- Identifier A can be used as a pointer to array element 0: Type T^*



Reference	Type	Value
val[4]	int	3
val	int *	x
val+1	int *	x+4
&val[2]	int *	x+8
val[5]	int	??
*(val+1)	int	5
val + i	int *	x+4i

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


1

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

4

Nested Array Example

```
#define PCOUNT 4
zip_dig pgh[PCOUNT] =
{ {1, 5, 2, 0, 6},
  {1, 5, 2, 1, 3},
  {1, 5, 2, 1, 7},
  {1, 5, 2, 2, 1} };
```




- “zip_dig pgh[4]” equivalent to “int pgh[4][5]”
 - Variable **pgh**: array of 4 elements, allocated contiguously
 - Each element is an array of 5 **int**'s, allocated contiguously
- “Row-Major” ordering of all elements in memory

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Nested Array Element Access



- Array Elements
 - $A[i][j]$ is element of type T , which requires K bytes
 - Address $A + i * (C * K) + j * K = A + (i * C + j) * K$



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Multidimensional (Nested) Arrays

- Declaration
 - $T \ A[R][C];$
 - 2D array of data type T
 - R rows, C columns
 - Type T element requires K bytes
- Array Size
 - $R * C * K$ bytes
- Arrangement
 - Row-Major Ordering



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Nested Array Row Access Code

```
int *get_pgh_zip(int index)
{
    return pgh[index];
}

# %rdi = index
leaq (%rdi,%rdi,4),%rax # 5 * index
leaq pgh(,%rax,4),%rax # pgh + (20 * index)
```

- Row Vector
 - $pgh[index]$ is array of 5 **int**'s
 - Starting address $pgh + 20 * index$
- Machine Code
 - Computes and returns address
 - Compute as $pgh + 4 * (index + 4 * index)$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Array Loop Example


```
void zinrc(zip_dig z) {
    size_t i;
    for (i = 0; i < ZLEN; i++)
        z[i]++;
}
```

```
# %rdi = z
movl $0, %eax # i = 0
jmp .L3 # goto middle
.L4: # loop:
addl $1, (%rdi,%rax,4) # z[i]++
addq $1, %rax # i++
.L3: # middle
cmpq $4, %rax # i:4
jbe .L4 # if <=, goto loop
ret
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Nested Array Row Access


- Row Vectors
 - $A[i]$ is array of C elements
 - Each element of type T requires K bytes
 - Starting address $A + i * (C * K)$



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Element Access in Multi-Level Array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



```
salq    $2, %rsi      # 4*digit
addq    univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax   # return *p
ret
```

Computation

- Element access **Mem[Mem[univ+8*index]+4*digit]**
- Must do two memory reads
 - First get pointer to row array
 - Then access element within array

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

15

16 X 16 Matrix Access

Array Elements

- Address $A + i * (C * K) + j * K$
- $C = 16, K = 4$

```
/* Get element a[i][j] */
int fix_ele(fix_matrix a, size_t i, size_t j) {
    return a[i][j];
}
```

```
# a in %rdi, i in %rsi, j in %rdx
salq    $6, %rsi      # 64*i
addq    %rsi, %rdi     # a + 64*i
movl    (%rdi,%rdx,4), %eax # M[a + 64*i + 4*j]
ret
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


16

Multi-Level Array Example

```
zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3
int *univ[UCOUNT] = {mit, cmu, ucb};
```

- Variable univ denotes array of 3 elements
- Each element is a pointer
 - 8 bytes
- Each pointer points to array of int's



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

14

N X N Matrix Code

Fixed dimensions

- Know value of N at compile time

Variable dimensions, explicit indexing

- Traditional way to implement dynamic arrays

Variable dimensions, implicit indexing

- Now supported by gcc

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele(fix_matrix a,
            size_t i, size_t j)
{
    return a[i][j];
}
```


```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element a[i][j] */
int vec_ele(size_t n, int *a,
            size_t i, size_t j)
{
    return a[IDX(n,i,j)];
}
```

```
/* Get element a[i][j] */
int var_ele(size_t n, int a[n][n],
            size_t i, size_t j) {
    return a[i][j];
}
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

17

Nested Array Element Access Code



```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

```
leaq    (%rdi,%rdi,4), %rax # 5*index
addl    %rax, %rsi         # 5*index+dig
movl    pgh(,%rsi,4), %eax # M[pgh + 4*(5*index+dig)]
```

Array Elements

- $pgh[index][dig]$ is int
- Address: $pgh + 20*index + 4*dig$
 - $= pgh + 4*(5*index + dig)$

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

13



Array Element Accesses

Nested array

```
int get_pgh_digit
(size_t index, size_t digit)
{
    return pgh[index][digit];
}
```

Multi-level array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



Accesses looks similar in C, but address computations very different:


Mem[pgh+20*index+4*digit] **Mem[Mem[univ+8*index]+4*digit]**

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

16

Structure Representation

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```




- Structure represented as block of memory
 - Big enough to hold all of the fields
- Fields ordered according to declaration
 - Even if another ordering could yield a more compact representation
- Compiler determines overall size + positions of fields
 - Machine-level program has no understanding of the structures in the source code

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

21

Structures & Alignment


Unaligned Data



```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

24

Today

- Arrays
 - One-dimensional
 - Multi-dimensional (nested)
 - Multi-level
- Structures
 - Allocation
 - Access
 - Alignment

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition


20

Following Linked List

C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    int i;
    struct rec *next;
};
```



Register	Value
%rdi	r
%rsi	val

```
.L11:                # loop:
    movslq    16(%rdi), %rax    # i = M[r+16]
    movl      %esi, (%rdi,%rax,4) # M[r+4*i] = val
    movq      24(%rdi), %rdi    # r = M[r+24]
    testq     %rdi, %rdi        # Test r
    jne       .L11              # if !=0 goto loop
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

23

n X n Matrix Access

- Array Elements
 - Address $A + i * (C * K) + j * K$
 - $C = n, K = 4$
 - Must perform integer multiplication

```
/* Get element a[i][j] */
int var_ele(size_t n, int a[n][n], size_t i, size_t j)
{
    return a[i][j];
}
```


```
# n in %rdi, a in %rsi, i in %rdx, j in %rcx
imulq    %rdx, %rdi    # n*i
leaq     (%rsi,%rdi,4), %rax # a + 4*n*i
movl     (%rax,%rcx,4), %eax # a + 4*n*i + 4*j
ret
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

19

Generating Pointer to Structure Member

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```



Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Compute as $r + 4 * idx$

```
int *get_ap
(struct rec *r, size_t idx)
{
    return &r->a[idx];
}
```

```
# r in %rdi, idx in %rsi
leaq     (%rdi,%rsi,4), %rax
ret
```


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

22

Satisfying Alignment with Structures

- Within structure:
 - Must satisfy each element's alignment requirement
- Overall structure placement
 - Each structure has alignment requirement K
 - K = Largest alignment of any element
 - Initial address & structure length must be multiples of K
- Example:
 - $K = 8$, due to `double` element

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```




Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

27

Accessing Array Elements

- Compute array offset $12 * \text{idx}$
 - `sizeof(S3)`, including alignment spacers
- Element j is at offset 8 within structure
- Assembler gives offset $a+8$
 - Resolved during linking

```
struct S3 {
    short i;
    float v;
    short j;
} a[10];
```



```
short get_j(int idx)
{
    return a[idx].j;
}
```

```
# %rdi = idx
leaq (%rdi,%rdi,2),%rax # 3*idx
movzwl a+8(%rax,4),%eax
```

Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

30

Specific Cases of Alignment (x86-64)

- 1 byte: `char`, ...
 - no restrictions on address
- 2 bytes: `short`, ...
 - lowest 1 bit of address must be 0_2
- 4 bytes: `int`, `float`, ...
 - lowest 2 bits of address must be 00_2
- 8 bytes: `double`, `long`, `char *`, ...
 - lowest 3 bits of address must be 000_2
- 16 bytes: `long double` (GCC on Linux)
 - lowest 4 bits of address must be 0000_2


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

28

Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element

```
struct S2 {
    double v;
    int i[2];
    char c;
} a[10];
```



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

29

Alignment Principles

- Aligned Data
 - Primitive data type requires K bytes
 - Address must be multiple of K
 - Required on some machines; advised on x86-64
- Motivation for Aligning Data
 - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans quad word boundaries
 - Virtual memory trickier when datum spans 2 pages
- Compiler
 - Inserts gaps in structure to ensure correct alignment of fields


Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

25

Meeting Overall Alignment Requirement

- For largest alignment requirement K
- Overall structure must be multiple of K

```
struct S2 {
    double v;
    int i[2];
    char c;
} *p;
```



Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

26

Saving Space

- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;  
→  
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- Effect (K=4)

c	3 bytes	i	d	3 bytes
---	---------	---	---	---------

i	c	d	2 bytes
---	---	---	---------

